

Sentence Segmentation Revisited: Is WtP The Game Changer?

Eric Asare, Mashrafi Monon, Fuad Thabit

New York University Abu Dhabi

{ea2525, mmm9886, ftt4473}@nyu.edu

Abstract

In this study, we conduct a primarily quantitative evaluation of the WtP Sentence Segmentation system, a tool reported for its advancements in handling multilingual and unpunctuated text in sentence segmentation tasks. Our investigation centers on the system’s performance across a range of text genres, primarily assessing its efficiency and effectiveness through metrics such as F1 score, precision, and recall. We specifically compare WtP’s performance with the Punkt model, a well-established tool in sentence segmentation, to understand how WtP fares against traditional approaches in this domain. Contrary to the system’s reported capabilities, our findings highlight some discrepancies in its performance when compared with the established Punkt model, particularly in the context of diverse linguistic scenarios. This paper details these comparative insights, focusing on the strengths and limitations of the WtP system, and discusses the implications of these findings for future enhancements in sentence segmentation methodologies within the field of Natural Language Processing (NLP).

1 Introduction

Sentence boundary detection (SBD) plays a crucial role in natural language processing (NLP), underpinning the success of various applications ranging from machine translation to information retrieval. Despite its significance, SBD has not been as extensively explored as other NLP tasks. This study focuses on the WtP Sentence Segmentation system, a novel approach that claims to significantly surpass existing sentence-segmentation tools. Employing a multilingual character-level language model, the WtP model represents a pioneering method in handling unpunctuated text across multiple languages.

However, in our attempt to corroborate the findings reported in the WtP paper, we encountered results that differed significantly. This paper aims

to detail these discrepancies and explore their implications in the field of NLP.

1.1 Motivation

The ability to accurately segment sentences is a fundamental requirement in NLP, yet it presents unique challenges, especially when dealing with languages that lack clear sentence-ending punctuation. With a plethora of tools available for sentence segmentation, the claim of WtP’s method to outperform existing systems by a significant margin necessitates an in-depth evaluation. Our study aims to explore the robustness of the WtP system and its potential limitations, contributing to a more comprehensive understanding of sentence segmentation effectiveness, especially in informal and unpunctuated text environments.

1.2 Study Objectives

Our study is guided by two primary objectives: firstly, to conduct a critical assessment of the WtP system’s performance across diverse linguistic scenarios; and secondly, to understand the broader implications of our findings for the field of sentence segmentation in NLP. We aim to provide a holistic perspective on the system’s strengths and potential areas for improvement. Through this approach, we intend to offer insights that could shape future research and development in sentence segmentation, guiding the advancement of more effective and versatile NLP tools.

1.3 Paper Organization

The remainder of this paper is organized as follows: Section 2 offers a review of related work in sentence boundary detection, providing context for our investigation. Section 3 details our methodology, including the evaluation metrics and datasets used to assess the WtP system. Section 4 presents our experimental results and analysis, shedding light on the performance of the WtP system in various

contexts. Finally, Section 5 discusses the implications of our findings for future research in sentence segmentation and concludes the paper with key takeaways from our study.

2 Related Work

Sentence segmentation in natural language processing (NLP) is essential yet challenging, particularly in multilingual contexts. The development of effective segmentation tools is a key area of research in NLP, involving a range of methodologies and approaches. This section explores the evolution of sentence boundary detection (SBD) techniques, focusing on both traditional and innovative methods, and contextualizes the development of the WtP Sentence Segmentation system within this landscape.

2.1 The Role of SBD in NLP

Read et al.'s work underscores the fundamental importance of SBD in NLP, highlighting the need for adaptable systems that can accommodate various text genres and languages without relying on text- or language-specific assumptions (Read et al., 2012). This broader view of SBD aligns with the principles underlying the WtP model, which employs self-supervision and multilingual capabilities, marking a departure from traditional SBD methods.

2.2 Evolution of Sentence Segmentation Approaches

Sentence segmentation approaches generally fall into three categories: rule-based, supervised machine learning, and unsupervised machine learning. Each category has its own strengths and weaknesses, and the choice of the most suitable approach often involves trade-offs, such as accuracy versus robustness to variations in text types.

2.2.1 Rule-Based and Unsupervised Learning

Punkt, an unsupervised learning model known for its abbreviation detection capability, exemplifies robustness across different text types (Kiss and Strunk, 2006). Its evaluation in eleven languages showcases its adaptability and effectiveness in handling linguistic variability. Punkt's approach, which relies on the collocation of words around periods and employs log-likelihood ratios, addresses ambiguities in sentence segmentation.

2.2.2 Supervised Learning and Advanced Techniques

The Elephant system represents a shift towards a hybrid approach, blending supervised and unsupervised learning for tokenization (Evang et al., 2013). This transition from rule-based to learning-based methods highlights the field's progression towards more dynamic and data-driven solutions.

2.2.3 Binary Classification Approach

Another novel approach in sentence segmentation is the binary classification model introduced by Wicks and Post. This method treats sentence segmentation as a binary classification problem, distinguishing between sentence-internal and sentence-ending positions. The core of their model is a two-layer Transformer, fine-tuned with parameters like vocabulary size, embedding size, and contextual window sizes.

Unique to their approach is the use of regular expressions to identify candidate sites for sentence boundary predictions, and the model's reliance on the context surrounding each token to make its determinations. Evaluations across 23 languages, encompassing 2,000 documents, have shown this approach to be highly effective, outperforming traditional rule-based systems like Punkt in multilingual contexts (Wicks and Post, 2021). This advancement highlights the potential of applying machine learning techniques, specifically binary classification, to the complex task of sentence segmentation in diverse linguistic scenarios.

2.3 Further Innovations in Sentence Segmentation

Recent advancements have introduced more sophisticated models like the bidirectional character-level language models (ChLMs). The WtP system, a prominent example of such innovation, adopts a self-supervised, multilingual, and punctuation-agnostic approach, marking a significant advance in the field (Minixhofer et al., 2023). Trained on the mC4 corpus in 85 languages, the WtP model's ability to predict sentence boundaries without relying on punctuation demonstrates its potential in handling diverse linguistic contexts, including those without standard sentence-ending punctuation.

2.4 Comparative Evaluation and Emerging Challenges

Comparative evaluations of new models like WtP against established tools such as Punkt are vital in

understanding their relative effectiveness. These analyses, which emphasize the need for comprehensive evaluation metrics, highlight the ongoing challenges in the field, such as dealing with ambiguity in punctuated contexts and the absence of standardized test sets and evaluation methods.

2.5 Implications for Future Research

The evolving landscape of sentence segmentation, characterized by models like WtP, suggests numerous avenues for future research. Enhancing the adaptability of segmentation models to various languages and text types, particularly in low-resource settings, remains a critical area of focus. This necessitates a balanced approach that addresses both accuracy and robustness in diverse linguistic environments.

The field of sentence segmentation is witnessing significant developments, with innovative models like WtP challenging traditional approaches and offering new possibilities. Future research should continue to explore and refine these models, ensuring their robustness and accuracy for real-world applications across a spectrum of languages and text types.

3 Methodology

Our methodology outlines the implementation and evaluation of the WtP Sentence Segmentation system, compared with the Punkt implementation, using a Google Colab GPU environment. This section details the steps taken to set up, execute, and evaluate the performance of these models.

Our decision to compare WtP with the Punkt model is further informed by the findings of Wicks and Post. Their research highlighted Punkt’s strong performance in English texts, achieving an impressive F1 score of 0.997, a testament to its efficacy in sentence segmentation tasks (Wicks and Post, 2021).

This performance, consistent with prior research, positions Punkt as a reliable benchmark for evaluating WtP’s capabilities, especially in handling English texts. The comparison aims to understand how WtP, with its advanced approaches, measures up against such a well-established standard in a diverse set of linguistic scenarios.

3.1 Implementation of WtP on Google Colab

The implementation process involved several key steps:

Environment Setup: The Python environment was set up in Google Colab, including the installation of the WtP package. This provided the necessary tools and libraries for running the WtP model.

Data Preparation: The MASC-3.0.0 dataset, containing texts from various genres, was prepared for processing. This dataset was chosen for its diversity, offering a wide range of linguistic contexts for robust testing.

Model Execution: The WtP model was employed to segment the text files into sentences. This process was crucial in understanding the model’s capability in handling different text genres and identifying sentence boundaries.

3.2 Dataset Source

For our study, the dataset utilized was the MASC-3.0.0 (Multi-Genre Corpus), which is a comprehensive and diverse dataset essential for our sentence segmentation analysis. This dataset can be accessed from the following website: <https://anc.org/data/masc/downloads/data-download/>. The inclusion of this dataset was pivotal in our methodology, as it provided a wide range of text genres, crucial for evaluating the sentence segmentation models under various linguistic contexts.

In addition to using the dataset in its original form, we also conducted a secondary analysis where newlines were removed from the raw files. This modification aimed to investigate the impact of newline characters on the segmentation performance of the WtP and Punkt systems. This step was crucial for determining how variations in data formatting could affect the accuracy of sentence segmentation models.

3.3 Additional Testing with OPUS Dataset

To further evaluate the WtP system’s performance and its adaptability to data used in its training process, we conducted tests using the OPUS dataset. This dataset was selected due to its use in WtP’s training, offering an opportunity to assess the system in a familiar context.

OPUS Dataset Preparation: The OPUS dataset, available at <https://data.statmt.org/opus-100-corpus/v1.0/supervised/bn-en/>, was used for this phase of testing. We focused on the development and test data within this dataset, which provided a diverse range of linguistic content for analysis.

Model Execution on OPUS: Unlike our initial methodology that compared WtP with Punkt, this phase solely assessed the performance of WtP on the OPUS dataset. This approach aimed to isolate WtP's behavior in a controlled environment, closely related to its training data.

Evaluation on OPUS: The performance of WtP on the OPUS dataset was evaluated under two conditions: with and without newline characters in the data. This was done to understand the impact of newline characters, a crucial aspect of WtP's training, on its segmentation capability.

Analysis of WtP's Adaptability: This analysis focused on examining how WtP's training on the OPUS data influenced its performance, particularly in handling newline characters and varying text lengths.

3.4 Evaluation Plan

The evaluation plan was designed to assess the effectiveness of the WtP system in comparison with the Punkt model. This involved:

Comparative Analysis with Punkt: Both the WtP and Punkt systems were used to process the same set of text files from the MASC-3.0.0 dataset. This direct comparison was aimed at highlighting the strengths and weaknesses of each system across various text genres.

Scoring and Metric Calculation: The performance of each system was evaluated based on precision, recall, and F-score metrics. These metrics were calculated by comparing the sentence boundary predictions made by each system against the ground truth annotations found in the '.xml' files of the dataset. The ground truth served as a benchmark for assessing the accuracy of the sentence segmentation.

Analysis Across Genres: The evaluation extended across different genres within the dataset, ensuring a comprehensive understanding of how each system performs under varying linguistic scenarios.

3.5 Scoring Process

The scoring process involved parsing the segmented outputs from both WtP and Punkt. For each text file, we compared the sentence boundary predictions against the ground truth. The ground truth was derived from the '.xml' files in the dataset, which contain accurate sentence boundary annotations.

3.6 Calculating Precision, Recall, and F-score

Precision was calculated as the ratio of correctly identified sentence boundaries to the total number of boundaries identified by the system. Recall measured the proportion of actual sentence boundaries correctly identified by the system. The F-score was computed as the harmonic mean of precision and recall, providing a balance between the two metrics.

3.7 Significance of the Evaluation

This methodology, enhanced by the additional newline removal test, provides deeper insights into the technical capabilities of the WtP system compared to Punkt. It contributes to a more comprehensive understanding of sentence segmentation in NLP. By testing these systems across various text genres and under different data formatting conditions, our study aims to uncover nuanced understandings of their applicability and potential areas for improvement in sentence segmentation technology.

4 Experimental Results

4.1 Performance Across Genres (With Newlines)

Our evaluation of the WtP and Punkt systems across various genres revealed notable differences in performance. The results are summarized as follows:

- In genres such as "journal" and "govt-docs", Punkt significantly outperformed WtP in precision, while WtP had higher recall.
- In "twitter" and "travel-guides", WtP showed a competitive F-score, indicating its effectiveness in certain contexts.
- Genres like "non-fiction", "fictions", and "blog" demonstrated the variability in effectiveness for both systems, with WtP showing higher recall but lower precision compared to Punkt.
- In more structured genres such as "technical" and "newspaper", Punkt generally had higher precision and F-score, though WtP maintained a high recall rate.
- For emails and letters, the performance varied, with WtP showing a notable precision in "letters/philanthropic-fundraising" but lower scores in other sub-genres.

The results are summarized in Table 2, which provides a comparative overview of the performance metrics for WtP and Punkt across various genres.

4.2 Performance Across Genres (With Newlines Removed)

In our extended evaluation, we conducted additional tests with newlines removed from the MASC-3.0.0 dataset. The results demonstrate how the removal of newlines impacts the performance of the WtP and Punkt systems:

- In "journal" and "govt-docs", WtP's precision increased significantly with newline removal, showing a more balanced performance compared to Punkt.
- For "twitter" and "travel-guides", WtP maintained competitive F-scores and displayed an improvement in precision.
- In genres like "non-fiction", "ficlets", and "blog", newline removal contributed to a slight increase in WtP's precision, balancing its performance with Punkt.
- In structured genres such as "technical" and "newspaper", WtP's precision and F-score improved, reducing the performance gap with Punkt.
- WtP's precision in "letters/philanthropic-fundraising" improved, reflecting better performance with newline removal.

These additional findings are summarized in Table 3, providing a comprehensive overview of the impact of newline removal on system performance.

4.3 Analysis of Precision, Recall, and F-score (With Newlines)

The detailed analysis of precision, recall, and F-score across genres highlights the strengths and weaknesses of each system:

- WtP excelled in recall across most genres, indicating its effectiveness in identifying sentence boundaries.
- Punkt showed superior precision in most genres, suggesting its accuracy in identifying correct sentence boundaries.

- The F-score, which balances precision and recall, varied across genres, with Punkt generally performing better. However, WtP showed competitive scores in certain contexts like "twitter" and "travel-guides".

4.4 Analysis of Precision, Recall, and F-score (With Newlines Removed)

The detailed analysis with newline removal further highlights the adaptability of each system:

- The removal of newlines generally improved WtP's precision across most genres, while its recall remained high.
- Punkt's performance remained relatively consistent, with minor variations in precision and recall across genres.
- The F-score comparison shows that newline removal can enhance WtP's overall effectiveness in certain genres.

4.5 WtP Performance on OPUS Dataset

Our exploration extended beyond the MASC-3.0.0 dataset to include focused evaluations of WtP using the OPUS dataset, a key component of its training data. This additional testing was pivotal in assessing WtP's performance in an environment mirroring its training context, particularly in terms of handling newline characters.

Consistency Across OPUS Datasets: WtP exhibited consistent performance across the development and test data within the OPUS dataset. This observation reinforces WtP's ability to adapt to data characteristics similar to those encountered during its training.

Performance with and without Newlines: Interestingly, WtP's performance improved on the OPUS dataset with the presence of newline characters, a factor contrasting with our findings from the MASC-3.0.0 dataset. When tested on the raw OPUS files without newline removal, WtP showed enhanced precision and recall, suggesting that its segmentation capabilities are closely aligned with the data formatting it was trained on.

The detailed performance metrics of WtP on the OPUS dataset are presented in Table 1. This table illustrates the system's precision, recall, and F-score under different conditions - with newlines and without newlines, across both development and test datasets. Notably, the presence of newlines in

the data led to higher precision and recall, underscoring the influence of familiarity of the nature of datasets used during training.

These findings not only highlight the impact of training data characteristics on the performance of sentence segmentation models like WtP but also emphasize the need for further research into how variations in data formatting and length affect segmentation accuracy.

4.6 Implications of the Results

The results from our tests, both with and without newlines, have significant implications for the application and development of sentence segmentation systems:

- **Impact of Data Formatting:** The notable differences in WtP’s performance when newlines are removed highlight the system’s sensitivity to data formatting. This sensitivity is likely due to WtP’s focus on newline prediction as a part of its segmentation process. In applications where data formatting varies or is inconsistent, WtP’s performance might fluctuate, indicating the need for preprocessing steps or system tuning to handle such variations.
- **Choice of System Based on Application Needs:** Depending on whether precision or recall is more critical for a specific application, the choice between WtP and Punkt may vary. For instance, in scenarios where missing a sentence boundary is particularly costly, WtP’s higher recall with newlines might be more beneficial, while Punkt’s consistent precision could be preferable in contexts requiring high accuracy in boundary detection.
- **Need for Genre-Specific Adaptations:** The variability in performance across different genres, observed in both sets of tests, underscores the necessity for genre-specific adaptations or training. Systems like WtP and Punkt might need to be fine-tuned or retrained on genre-specific data to optimize their performance.
- **Consideration of Training Data Characteristics:** The differences observed in WtP’s performance with and without newlines suggest that the characteristics of the training data play a crucial role. If the training data predominantly included or excluded newlines, this

Condition	Precision	Recall	F-score
With New Lines (Dev)	0.633	0.986	0.771
With New Lines (Test)	0.653	0.981	0.784
Without Newlines (Dev)	0.643	0.883	0.744
Without Newlines (Test)	0.638	0.8915	0.7435

Table 1: WtP Performance on OPUS Dataset.

might have influenced the system’s learning process, thereby impacting its performance in real-world applications.

These insights underline the importance of not only the technological aspects of sentence segmentation systems but also the characteristics of the data they are trained and tested on. Understanding these dynamics is crucial for the effective deployment and further development of such systems in diverse NLP applications.

5 Discussion and Conclusion

Our comprehensive study of the WtP and Punkt sentence segmentation systems, including tests with and without newline characters, reveals intricate complexities in achieving effective sentence segmentation. This research underscores the critical importance of considering both the nature of the text, specific application requirements, and data formatting when deploying sentence segmentation models in NLP.

5.1 Insights and Implications

The comparative analysis between WtP and Punkt, enriched by the newline removal tests, has provided valuable insights into their respective strengths, weaknesses, and adaptability to data formatting. WtP’s novel approach, utilizing bidirectional character-level language models, demonstrates significant potential and a notable sensitivity to newline characters. This sensitivity, evident in its varying performance across different genres, underscores the influence of its training on data characteristics and highlights the importance of preprocessing in certain contexts. Particularly, WtP’s granular analysis seems to excel in genres where contextual nuances and character-level subtleties are paramount, such as informal texts or languages lacking explicit sentence-ending punctuation.

In contrast, Punkt’s rule-based approach, leveraging the collocation of words around periods and log-likelihood ratios, shows a robust consistency in more structured texts like formal documents or

news articles, where punctuation and word patterns are more predictable. This distinction in methodologies and their corresponding effectiveness across genres underscores the necessity of choosing segmentation models that align with the specific requirements of the task at hand. The balance between precision and recall in different text types is a critical consideration, as evidenced by our findings.

These revelations provide valuable guidance for future developments in NLP, emphasizing the need to tailor sentence segmentation tools to the specific characteristics and demands of the text genre. As the field evolves, understanding these dynamics will be crucial for optimizing the effectiveness of NLP applications.

5.2 Future Work

Given the differences observed between our results and those reported in the original WtP paper, a potential future step could involve reaching out to the authors of the WtP study. Such an engagement could offer insights into these discrepancies and might guide subsequent experimental approaches or revisions to the WtP system, contributing to the broader discourse in NLP technology development.

Moreover, an interesting avenue for research could be the development of an ensemble model that combines the strengths of WtP and Punkt. Considering the distinct advantages each system demonstrated in our study, a model that strategically merges the high precision of Punkt in certain genres with the effective sentence boundary predictions of WtP, particularly in cases where its precision exceeds 50%, could potentially yield a more adaptable and robust tool for sentence segmentation. Exploring this hypothesis through targeted experiments could provide valuable contributions to the field, suggesting novel methods for optimizing sentence segmentation in various linguistic contexts.

5.3 Conclusion and Key Takeaways

Our study, enriched by the additional newline character tests, demonstrates the crucial importance of evaluating sentence segmentation models using a diverse array of text genres, robust metrics, and considering data formatting. The insights gained contribute significantly to the ongoing discourse in sentence segmentation efficacy and pave the way for future advancements in this critical area of NLP.

The key takeaway is that the choice of a segmentation model should be informed by the text genre, application requirements, and data characteristics. As the field of NLP evolves, our approaches to sentence segmentation must also adapt, ensuring they are precise, robust, and versatile enough to meet the diverse challenges presented by different text genres, linguistic contexts, and data formats.

References

- Kilian Evang, Valerio Basile, Grzegorz Chrupała, and Johan Bos. 2013. [Elephant: Sequence labeling for word and sentence segmentation](#). In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1422–1426, Seattle, Washington, USA. Association for Computational Linguistics.
- Tibor Kiss and Jan Strunk. 2006. [Unsupervised multilingual sentence boundary detection](#). *Computational Linguistics*, 32(4):485–525.
- Benjamin Minixhofer, Jonas Pfeiffer, and Ivan Vulić. 2023. [Where’s the point? self-supervised multilingual punctuation-agnostic sentence segmentation](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7215–7235, Toronto, Canada. Association for Computational Linguistics.
- Jonathon Read, Rebecca Dridan, Stephan Oepen, and Lars Jørgen Solberg. 2012. [Sentence boundary detection: A long solved problem?](#) In *Proceedings of COLING 2012: Posters*, pages 985–994, Mumbai, India. The COLING 2012 Organizing Committee.
- Rachel Wicks and Matt Post. 2021. [A unified approach to sentence segmentation of punctuated text in many languages](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 3995–4007, Online. Association for Computational Linguistics.

Table 2: Performance Comparison of WtP and Punkt Across Genres (With Newlines)

Genre	WtP			Punkt		
	Precision	Recall	F-score	Precision	Recall	F-score
Journal	0.083	0.969	0.153	0.909	0.831	0.868
Govt-Docs	0.086	0.825	0.156	0.862	0.903	0.882
Essays	0.264	0.806	0.398	0.763	0.818	0.790
Twitter	0.470	0.842	0.604	0.931	0.637	0.757
Travel Guides	0.673	0.965	0.793	0.934	0.959	0.946
Non-fiction	0.742	0.754	0.748	0.849	0.904	0.876
Ficlets	0.588	0.669	0.626	0.910	0.668	0.771
Blog	0.137	0.844	0.235	0.910	0.826	0.866
Movie Script	0.296	0.607	0.398	0.873	0.694	0.774
Fiction	0.490	0.655	0.560	0.962	0.623	0.756
Jokes	0.765	0.574	0.656	0.816	0.629	0.711
Technical	0.451	0.930	0.607	0.932	0.898	0.915
Letters/Solicitation Brochures	0.349	0.748	0.476	0.928	0.829	0.876
Letters/Philanthropic Fundraising	0.734	0.853	0.789	0.943	0.792	0.861
Newspaper/WSJ	0.643	0.613	0.628	0.784	0.908	0.841
Newspaper/NYT	0.102	0.772	0.179	0.885	0.783	0.830
Newspaper/Unknown	0.500	0.500	0.500	0.929	0.722	0.812
Email/W3C	0.201	0.859	0.326	0.885	0.672	0.764
Email/Spam	0.193	0.824	0.313	0.848	0.693	0.763
Email/Enron	0.147	0.836	0.250	0.897	0.694	0.782

Table 3: Performance Comparison of WtP and Punkt Across Genres (With Newlines Removed)

Genre	WtP (No Newlines)			Punkt (No Newlines)		
	Precision	Recall	F-score	Precision	Recall	F-score
Journal	0.7158	0.9807	0.8276	0.9088	0.8310	0.8682
Govt-Docs	0.6923	0.9439	0.7987	0.8619	0.9029	0.8820
Essays	0.4829	0.8975	0.6279	0.7633	0.8182	0.7898
Twitter	0.5042	0.8816	0.6415	0.9313	0.6373	0.7567
Travel Guides	0.8021	0.9692	0.8778	0.9337	0.9587	0.9460
Non-fiction	0.7929	0.9457	0.8626	0.8493	0.9041	0.8758
Ficlets	0.6385	0.7467	0.6884	0.9099	0.6684	0.7707
Blog	0.6975	0.9203	0.7936	0.9103	0.8265	0.8664
Movie Script	0.5498	0.6907	0.6123	0.8731	0.6945	0.7736
Fiction	0.8890	0.6976	0.7817	0.9620	0.6227	0.7561
Jokes	0.7862	0.6897	0.7348	0.8162	0.6294	0.7107
Technical	0.7235	0.8990	0.8018	0.9322	0.8978	0.9147
Letters/Solicitation Brochures	0.6037	0.9296	0.7320	0.9285	0.8292	0.8760
Letters/Philanthropic Fundraising	0.7213	0.8980	0.8000	0.9432	0.7921	0.8611
Newspaper/WSJ	0.7129	0.9057	0.7978	0.7837	0.9082	0.8414
Newspaper/NYT	0.6262	0.8214	0.7106	0.8847	0.7825	0.8305
Newspaper/Unknown	0.3902	0.8889	0.5424	0.9286	0.7222	0.8125
Email/W3C	0.4824	0.7756	0.5948	0.8847	0.6720	0.7638
Email/Spam	0.3886	0.9061	0.5439	0.8481	0.6933	0.7630
Email/Enron	0.4078	0.7666	0.5323	0.8970	0.6936	0.7823